

- A stack is a useful tool both for translating an infix to a postfix expression and for evaluating a postfix expression.

Chapter 5: Linked Lists

Overview

In [Chapter 2, "Arrays,"](#) we saw that arrays had certain disadvantages as data storage structures. In an unordered array, searching is slow, whereas in an ordered array, insertion is slow. In both kinds of arrays deletion is slow. Also, the size of an array can't be changed after it's created.

In this chapter we'll look at a data storage structure that solves some of these problems: the *linked list*. Linked lists are probably the second most commonly used general-purpose storage structures after arrays.

The linked list is a versatile mechanism suitable for use in many kinds of general-purpose databases. It can also replace an array as the basis for other storage structures such as stacks and queues. In fact, you can use a linked list in many cases where you use an array (unless you need frequent random access to individual items using an index).

Linked lists aren't the solution to all data storage problems, but they are surprisingly versatile and conceptually simpler than some other popular structures such as trees. We'll investigate their strengths and weaknesses as we go along.

In this chapter we'll look at simple linked lists, double-ended lists, sorted lists, doubly linked lists, and lists with iterators (an approach to random access to list elements). We'll also examine the idea of Abstract Data Types (ADTs) and see how stacks and queues can be viewed as ADTs and how they can be implemented as linked lists instead of arrays.

Links

In a linked list, each data item is embedded in a *link*. A link is an object of a class called something like `Link`. Because there are many similar links in a list, it makes sense to use a separate class for them, distinct from the linked list itself. Each link object contains a reference (usually called `next`) to the next link in the list. A field in the list itself contains a reference to the first link. This is shown in [Figure 5.1](#).

Here's part of the definition of a class `Link`. It contains some data and a reference to the next link.

```
class Link
{
    public int iData;      // data
    public double dData;  // data
    public Link next;     // reference to next link
}
```

This kind of class definition is sometimes called *self-referential* because it contains a field—called `next` in this case—of the same type as itself.

We show only two data items in the link: an `int` and a `double`. In a typical application there would be many more. A personnel record, for example, might have name, address, Social Security number, title, salary, and many other fields. Often an object of a class that contains this data is used instead of the items:

```

class Link
{
    public inventoryItem iI; // object holding data
    public Link next;       // reference to next link
}

```



Figure 5.1: Links in a list

References and Basic Types

It's easy to get confused about references in the context of linked lists, so let's review how they work.

It may seem odd that you can put a field of type `Link` inside the class definition of this same type. Wouldn't the compiler be confused? How can it figure out how big to make a `Link` object if a link contains a link and the compiler doesn't already know how big a `Link` object is?

The answer is that a `Link` object doesn't really contain another `Link` object, although it may look like it does. The `next` field of type `Link` is only a *reference* to another link, not an object.

A reference is a number that *refers to* an object. It's the object's address in the computer's memory, but you don't need to know its value; you just treat it as a magic number that tells you where the object is. In a given computer/operating system, all references, no matter what they refer to, are the same size. Thus it's no problem for the compiler to figure out how big this field should be, and thereby construct an entire `Link` object.

Note that in Java, primitive types like `int` and `double` are stored quite differently than objects. Fields containing primitive types do not contain references, but actual numerical values like 7 or 3.14159. A variable definition like

```
double salary = 65000.00;
```

creates a space in memory and puts the number 65000.00 into this space. However, a reference to an object, like

```
Link aLink = someLink;
```

puts a reference to an object of type `Link`, called `someLink`, into `aLink`. The `someLink` object isn't moved, or even created, by this statement; it must have been created before. To create an object you must always use `new`:

```
Link someLink = new Link();
```

Even the `someLink` field doesn't hold an object, it's still just a reference. The object is somewhere else in memory, as shown in Figure 5.2.

Other languages, like C++, handle objects quite differently than Java. In C++ a field like

```
Link next;
```

actually contains an object of type `Link`. You can't write a self-referential class definition in C++ (although you can put a pointer to a `Link` in class `Link`; a pointer is similar to a reference). C++ programmers should keep in mind how Java handles objects; it may be counter intuitive.




Figure 5.2: Objects and references in memory

Relationship, Not Position

Let's examine one of the major ways in which linked lists differ from arrays. In an array each item occupies a particular position. This position can be directly accessed using an index number. It's like a row of houses: you can find a particular house using its address.

In a list the only way to find a particular element is to follow along the chain of elements. It's more like human relations. Maybe you ask Harry where Bob is. Harry doesn't know, but he thinks Jane might know, so you go and ask Jane. Jane saw Bob leave the office with Sally, so you call Sally's cell phone. She dropped Bob off at Peter's office, so...but you get the idea. You can't access a data item directly; you must use relationships between the items to locate it. You start with the first item, go to the second, then the third, until you find what you're looking for.

The LinkedList Workshop Applet

The LinkedList Workshop applet provides three list operations. You can insert a new data item, search for a data item with a specified key, and delete a data item with a specified key. These operations are the same ones we explored in the Array Workshop applet in [Chapter 2](#); they're suitable for a general-purpose database application.

Figure 5.3 shows how the LinkedList Workshop applet looks when it's started up. Initially there are 13 links on the list.



Figure 5.3: The LinkList Workshop applet

Insert

If you think 13 is an unlucky number, you can insert a new link. Click on the Ins button, and you'll be prompted to enter a key value between 0 and 999. Subsequent presses will generate a link with this data in it, as shown in Figure 5.4.

In this version of a linked list, new links are always inserted at the beginning of the list. This is the simplest approach, although it's also possible to insert links anywhere in the list, as we'll see later.

A final press on Ins will redraw the list so the newly inserted link lines up with the other links. This redrawing doesn't represent anything happening in the program itself, it just makes the display look neater.

Find

The Find button allows you find a link with a specified key value. When prompted, type in the value of an existing link, preferably one somewhere in the middle of the list. As you continue to press the button, you'll see the red arrow move along the list, looking for the link. A message informs you when it finds it. If you type a nonexistent key value, the arrow will search all the way to the end of the list before reporting that the item can't be found.



Figure 5.4: A new link being inserted

Delete

You can also delete a key with a specified value. Type in the value of an existing link, and

repeatedly press Del. Again the arrow will move along the list, looking for the link. When it finds it, it simply removes it and connects the arrow from the previous link straight across to the following link. This is how links are removed: the reference to the preceding link is changed to point to the following link.

A final keypress redraws the picture, but again this just provides evenly spaced links for aesthetic reasons; the length of the arrows doesn't correspond to anything in the program.

Unsorted and Sorted

The LinkList Workshop applet can create both unsorted and sorted lists. Unsorted is the default. We'll show how to use the applet for sorted lists when we discuss them later in this chapter.

A Simple Linked List

Our first example program, `linkList.java`, demonstrates a simple linked list. The only operations allowed in this version of a list are

- Inserting an item at the beginning of the list
- Deleting the item at the beginning of the list
- Iterating through the list to display its contents

These operations are fairly easy to carry out, so we'll start with them. (As we'll see later, these operations are also all you need to use a linked list as the basis for a stack.)

Before we get to the complete `linkList.java` program, we'll look at some important parts of the `Link` and `LinkList` classes.

The Link Class

You've already seen the data part of the `Link` class. Here's the complete class definition:

```
class Link
{
    public int iData;           // data item
    public double dData;       // data item
    public Link next;          // next link in list

// -----
-
    public Link(int id, double dd) // constructor
    {
        iData = id;               // initialize data
        dData = dd;               // ('next' is automatically
    }                             // set to null)

// -----
-
    public void displayLink()      // display ourself
    {
        System.out.print("{ " + iData + ", " + dData + " } ");
    }
}
```

```

    }

} // end class Link

```

In addition to the data, there's a constructor and a method, `displayLink()`, that displays the link's data in the format `{22, 33.9}`. Object purists would probably object to naming this method `displayLink()`, arguing that it should be simply `display()`. This would be in the spirit of polymorphism, but it makes the listing somewhat harder to understand when you see a statement like

```
current.display();
```

and you've forgotten whether `current` is a `Link` object, a `LinkedList` object, or something else.

The constructor initializes the data. There's no need to initialize the `next` field, because it's automatically set to `null` when it's created. (Although it could be set to `null` explicitly, for clarity.) The `null` value means it doesn't refer to anything, which is the situation until the link is connected to other links.

We've made the storage type of the `Link` fields (`iData` and so on) `public`. If they were `private` we would need to provide public methods to access them, which would require extra code, thus making the listing longer and harder to read. Ideally, for security we would probably want to restrict `Link`-object access to methods of the `LinkedList` class. However, without an inheritance relationship between these classes, that's not very convenient. We could use the default access specifier (no keyword) to give the data *package access* (access restricted to classes in the same directory) but that has no effect in these demo programs, which only occupy one directory anyway. The `public` specifier at least makes it clear that this data isn't private.

The LinkedList Class

The `LinkedList` class contains only one data item: a reference to the first link on the list. This reference is called `first`. It's the only permanent information the list maintains about the location of any of the links. It finds the other links by following the chain of references from `first`, using each link's `next` field.

```

class LinkedList
{
    private Link first;           // ref to first link on list

// -----
-
    public void LinkedList()      // constructor
    {
        first = null;           // no items on list yet
    }

// -----
-
    public boolean isEmpty()      // true if list is empty
    {
        return (first==null);
    }
}

```



```
// -----
-
    ...                               // other methods go here

}
```

The constructor for `LinkedList` sets `first` to `null`. This isn't really necessary because as we noted, references are set to `null` automatically when they're created. However, the explicit constructor makes it clear that this is how `first` begins.

When `first` has the value `null`, we know there are no items on the list. If there were any items, `first` would contain a reference to the first one. The `isEmpty()` method uses this fact to determine whether the list is empty.

The `insertFirst()` Method

The `insertFirst()` method of `LinkedList` inserts a new link at the beginning of the list. This is the easiest place to insert a link, because `first` already points to the first link. To insert the new link, we need only set the `next` field in the newly created link to point to the old first link, and then change `first` so it points to the newly created link. This is shown in [Figure 5.5](#).

In `insertFirst()` we begin by creating the new link using the data passed as arguments. Then we change the link references as we just noted.

```
                                // insert at start of list
public void insertFirst(int id, double dd)
{
    Link newLink = new Link(id, dd);    // make new link
    newLink.next = first;                // newLink --> old first
    first = newLink;                    // first --> newLink
}
```

The arrows `-->` in the comments in the last two statements mean that a link (or the `first` field) connects to the next (downstream) link. (In doubly linked lists we'll see upstream connections as well, symbolized by `<--` arrows.) Compare these two statements with [Figure 5.5](#). Make sure you understand how the statements cause the links to be changed, as shown in the figure. This kind of reference-manipulation is the heart of linked list algorithms.

The `deleteFirst()` Method

The `deleteFirst()` method is the reverse of `insertFirst()`. It disconnects the first link by rerouting `first` to point to the second link. This second link is found by looking at the `next` field in the first link.

```
public Link deleteFirst()        // delete first item
{
    Link temp = first;           // (assumes list not empty)
    first = first.next;          // save reference to link
    return temp;                 // delete it: first-->old next
}
```




Figure 5.5: Inserting a new link

The second statement is all you need to remove the first link from the list. We choose to also return the link, for the convenience of the user of the linked list, so we save it in `temp` before deleting it, and return the value of `temp`. Figure 5.6 shows how `first` is rerouted to delete the object.

In C++ and similar languages, you would need to worry about deleting the link itself after it was disconnected from the list. It's in memory somewhere, but now nothing refers to it. What will become of it? In Java, the garbage collection process will destroy it at some point in the future; it's not your responsibility.




Figure 5.6: Deleting a link

Notice that the `deleteFirst()` method assumes the list is not empty. Before calling it, your program should verify this with the `isEmpty()` method.

The `displayList()` Method

To display the list, you start at `first` and follow the chain of references from link to link. A variable `current` points to (or technically *refers to*) each link in turn. It starts off pointing to `first`, which holds a reference to the first link. The statement

```
current = current.next;
```


changes `current` to point to the next link, because that's what's in the `next` field in each link. Here's the entire `displayList()` method:

```
public void displayList()
{
    System.out.print("List (first-->last): ");
    Link current = first;          // start at beginning of list
    while(current != null)         // until end of list,
    {
        current.displayLink();    // print data
        current = current.next;   // move to next link
    }
    System.out.println("");
}
```

The end of the list is indicated by the `next` field in the last link pointing to `null` rather than another link. How did this field get to be `null`? It started that way when the link was created and was never given any other value because it was always at the end of the list. The `while` loop uses this condition to terminate itself when it reaches the end of the list. Figure 5.7 shows how `current` steps along the list.

At each link, the `displayList()` method calls the `displayLink()` method to display the data in the link.

The `linkList.java` Program

Listing 5.1 shows the complete `linkList.java` program. You've already seen all the components except the `main()` routine.




Figure 5.7: Stepping along the list

Listing 5.1 The `linkList.java` Program

```
// linkList.java
// demonstrates linked list
// to run this program: C>java LinkListApp
////////////////////////////////////
class Link
```



```

        public Link deleteFirst()          // delete first item
        {
            Link temp = first;             // (assumes list not empty)
            first = first.next;             // save reference to link
            first = first.next;             // delete it: first-->old
next
            return temp;                    // return deleted link
        }

// -----
-
    public void displayList()
    {
        System.out.print("List (first-->last): ");
        Link current = first;              // start at beginning of list
        while(current != null)              // until end of list,
        {
            current.displayLink();          // print data
            current = current.next;          // move to next link
        }
        System.out.println("");
    }

// -----
-
    } // end class LinkList

////////////////////////////////////

class LinkListApp
{
    public static void main(String[] args)
    {
        LinkList theList = new LinkList(); // make new list

        theList.insertFirst(22, 2.99);      // insert four items
        theList.insertFirst(44, 4.99);
        theList.insertFirst(66, 6.99);
        theList.insertFirst(88, 8.99);

        theList.displayList();               // display list

        while( !theList.isEmpty() )          // until it's empty,
        {
            Link aLink = theList.deleteFirst(); // delete link
            System.out.print("Deleted ");      // display it
            aLink.displayLink();
            System.out.println("");
        }
        theList.displayList();               // display list
    } // end main()

} // end class LinkListApp

```

In `main()` we create a new list, insert four new links into it with `insertFirst()`, and display it. Then, in the `while` loop, we remove the items one by one with `deleteFirst()` until the list is empty. The empty list is then displayed. Here's the output from `linkList.java`:

```
List (first-->last): {88, 8.99} {66, 6.99} {44, 4.99} {22,
2.99}
Deleted {88, 8.99}
Deleted {66, 6.99}
Deleted {44, 4.99}
Deleted {22, 2.99}
List (first-->last):
```

Finding and Deleting Specified Links

Our next example program adds methods to search a linked list for a data item with a specified key value, and to delete an item with a specified key value. These, along with insertion at the start of the list, are the same operations carried out by the `LinkList` Workshop applet. The complete `linkList2.java` program is shown in Listing 5.2.

Listing 5.2 The `linkList2.java` Program

```
// linkList2.java
// demonstrates linked list
// to run this program: C>java LinkList2App
/////////////////////////////////////////////////////////////////
class Link
{
    public int iData;           // data item (key)
    public double dData;       // data item
    public Link next;          // next link in list

// -----
-
    public Link(int id, double dd) // constructor
    {
        iData = id;
        dData = dd;
    }

// -----
-
    public void displayLink()      // display ourself
    {
        System.out.print("{ " + iData + ", " + dData + " } ");
    }
} // end class Link

/////////////////////////////////////////////////////////////////

class LinkList
{
    private Link first;           // ref to first link on list
```

```

// -----
-
    public LinkedList()           // constructor
    {
        first = null;           // no links on list yet
    }

// -----
-
    public void insertFirst(int id, double dd)
    {
        // make new link
        Link newLink = new Link(id, dd);
        newLink.next = first;    // it points to old first
link    first = newLink;        // now first points to this
    }

// -----
-
    public Link find(int key)     // find link with given key
    {
        // (assumes non-empty list)
        Link current = first;    // start at 'first'
        while(current.iData != key) // while no match,
        {
            if(current.next == null) // if end of list,
                return null;         // didn't find it
            else                     // not end of list,
                current = current.next; // go to next link
        }
        return current;          // found it
    }

// -----
-
    public Link delete(int key)   // delete link with given key
    {
        // (assumes non-empty list)
        Link current = first;    // search for link
        Link previous = first;
        while(current.iData != key)
        {
            if(current.next == null)
                return null;      // didn't find it
            else
            {
                previous = current; // go to next link
                current = current.next;
            }
        }
        // found it
        if(current == first)
            first = first.next;   // change first
        else
            previous.next = current.next; // bypass it
        return current;
    }

```

```

    }

// -----
-
    public void displayList()      // display the list
    {
        System.out.print("List (first-->last): ");
        Link current = first;      // start at beginning of list
        while(current != null)      // until end of list,
        {
            current.displayLink();  // print data
            current = current.next;  // move to next link
        }
        System.out.println("");
    }

// -----
-
    } // end class LinkList

////////////////////////////////////

class LinkList2App
{
    public static void main(String[] args)
    {
        LinkList theList = new LinkList(); // make list

        theList.insertFirst(22, 2.99);      // insert 4 items
        theList.insertFirst(44, 4.99);
        theList.insertFirst(66, 6.99);
        theList.insertFirst(88, 8.99);

        theList.displayList();               // display list

        Link f = theList.find(44);           // find item
        if( f != null)
            System.out.println("Found link with key " + f.iData);
        else
            System.out.println("Can't find link");

        Link d = theList.delete(66);          // delete item
        if( d != null )
            System.out.println("Deleted link with key " +
d.iData);
        else
            System.out.println("Can't delete link");
        theList.displayList();               // display list
    } // end main()

} // end class LinkList2App

```

The `main()` routine makes a list, inserts four items, and displays the resulting list. It then searches for the item with key 44, deletes the item with key 66, and displays the list

again. Here's the output:

```
List (first-->last): {88, 8.99} {66, 6.99} {44, 4.99} {22, 2.99}
Found link with key 44
Deleted link with key 66
List (first-->last): {88, 8.99} {44, 4.99} {22, 2.99}
```

The find() Method

The `find()` method works much like the `displayList()` method seen in the last program. The reference `current` initially points to `first`, and then steps its way along the links by setting itself repeatedly to `current.next`. At each link, `find()` checks if that link's key is the one it's looking for. If it is, it returns with a reference to that link. If it reaches the end of the list without finding the desired link, it returns `null`.

The delete() Method

The `delete()` method is similar to `find()` in the way it searches for the link to be deleted. However, it needs to maintain a reference not only to the current link (`current`), but to the link preceding the current link (`previous`). This is because, if it deletes the current link, it must connect the preceding link to the following link, as shown in [Figure 5.8](#). The only way to tell where the preceding link is, is to maintain a reference to it.

At each cycle through the `while` loop, just before `current` is set to `current.next`, `previous` is set to `current`. This keeps it pointing at the link preceding `current`.

To delete the current link once it's found, the `next` field of the previous link is set to the next link. A special case arises if the current link is the first link because the first link is pointed to by the `LinkedList`'s `first` field and not by another link. In this case the link is deleted by changing `first` to point to `first.next`, as we saw in the last program with the `deleteFirst()` method. Here's the code that covers these two possibilities:

```
                                // found it

if(current == first)           // if first link,
    first = first.next;        //   change first

else                            // otherwise,

    previous.next = current.next; //   bypass link
```

Other Methods

We've seen methods to insert and delete items at the start of a list and to find a specified item and delete a specified item. You can imagine other useful list methods. For example, an `insertAfter()` method could find a link with a specified key value and insert a new link following it. We'll see such a method when we talk about list iterators at the end of this chapter.




Figure 5.8: Deleting a specified link

Double-Ended Lists

A double-ended list is similar to an ordinary linked list, but it has one additional feature: a reference to the last link as well as to the first. Figure 5.9 shows what this looks like.



Figure 5.9: A double-ended list

The reference to the last link permits you to insert a new link directly at the end of the list as well as at the beginning. Of course you can insert a new link at the end of an ordinary single-ended list by iterating through the entire list until you reach the end, but this is inefficient.

Access to the end of the list as well as the beginning makes the double-ended list suitable for certain situations that a single-ended list can't handle efficiently. One such situation is implementing a queue; we'll see how this works in the next section.

Listing 5.3 contains the `firstLastList.java` program, which demonstrates a double-ended list. (Incidentally, don't confuse the double-ended list with the doubly linked list, which we'll explore later in this chapter.)

Listing 5.3 The `firstLastList.java` Program

```
// firstLastList.java
// demonstrates list with first and last references
// to run this program: C>java FirstLastApp
////////////////////////////////////
class Link
{
    public double dData;           // data item
    public Link next;              // next link in list

    // -----
    -
```